



Privacy-Preserving Machine Learning in TensorFlow

Jason Mancuso
Yann Dupis



DropoutLabs



Overview

Privacy-preserving machine learning and why we may need it

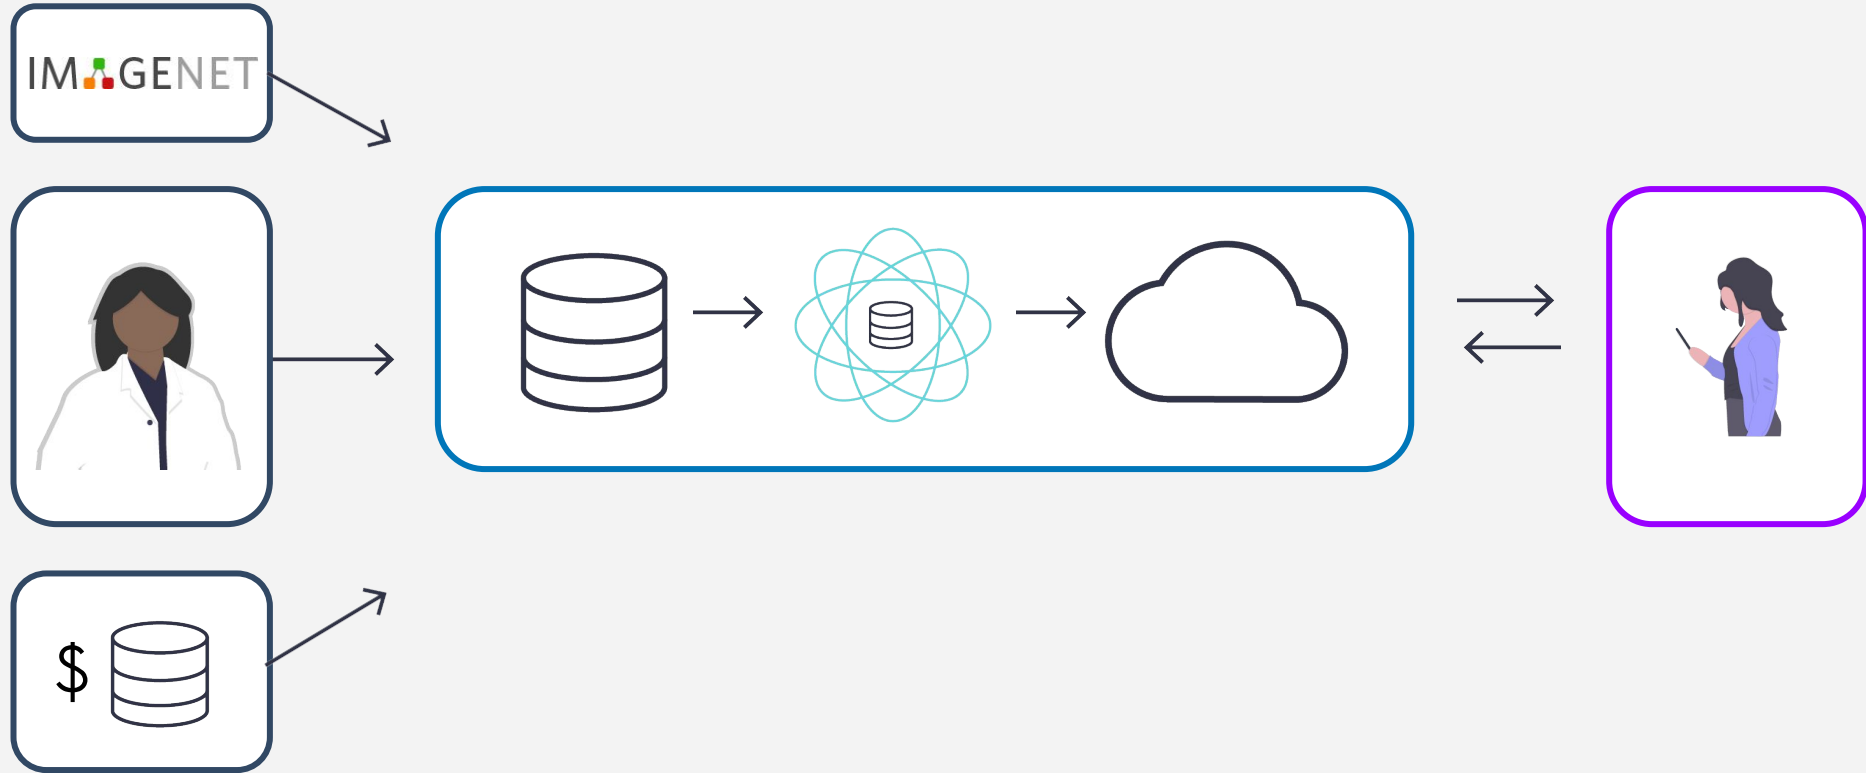
- Training **Remote execution** using PySyft-TensorFlow
- Training with **Differential Privacy**
- **Federated Learning** with secure aggregation using TF Encrypted
- **Encrypted Predictions** using TF Encrypted

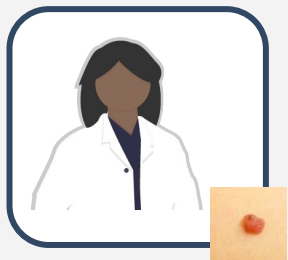
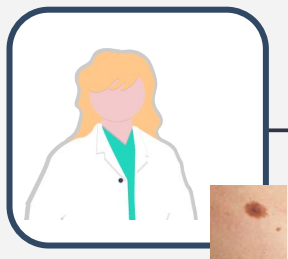
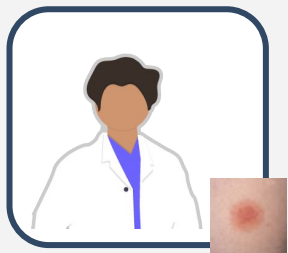
Why?

Privacy in machine learning

Privacy leakage creates
bottlenecks

Machine Learning Workflow



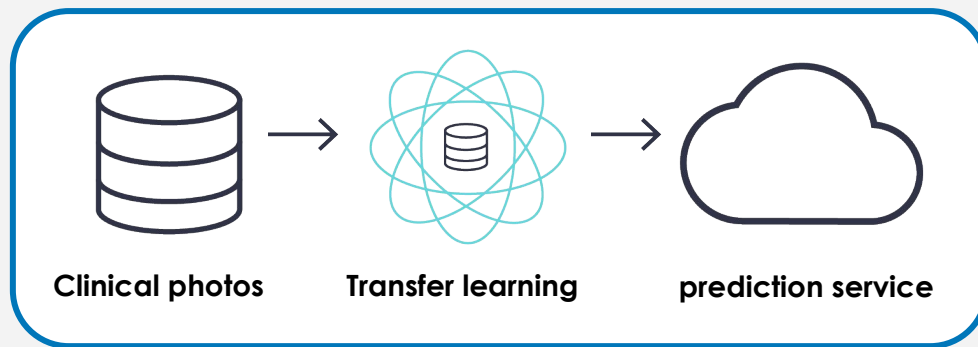


Skin Cancer Image Classification

Brett Kuprel

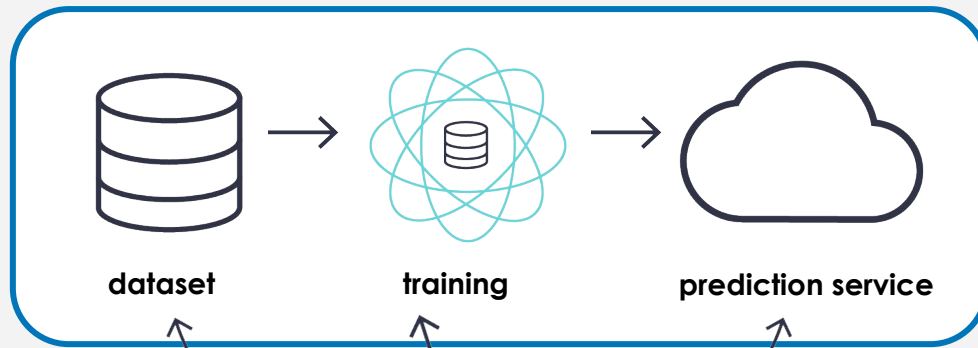
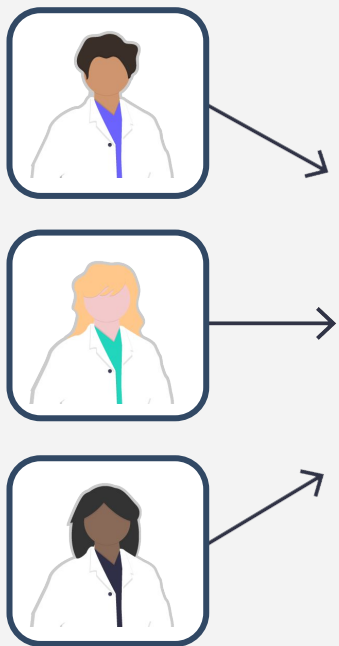
12:30-12:40pm

Join Brett Kuprel, and see how TensorFlow was used by the artificial intelligence lab and medical school of Stanford to classify skin cancer images. He'll describe the project steps: from acquiring a dataset, training a deep network, and evaluating the results. To wrap up, Brett will give his take on the future of skin cancer image classification.



Bottlenecks

data access (liability and controlled use)



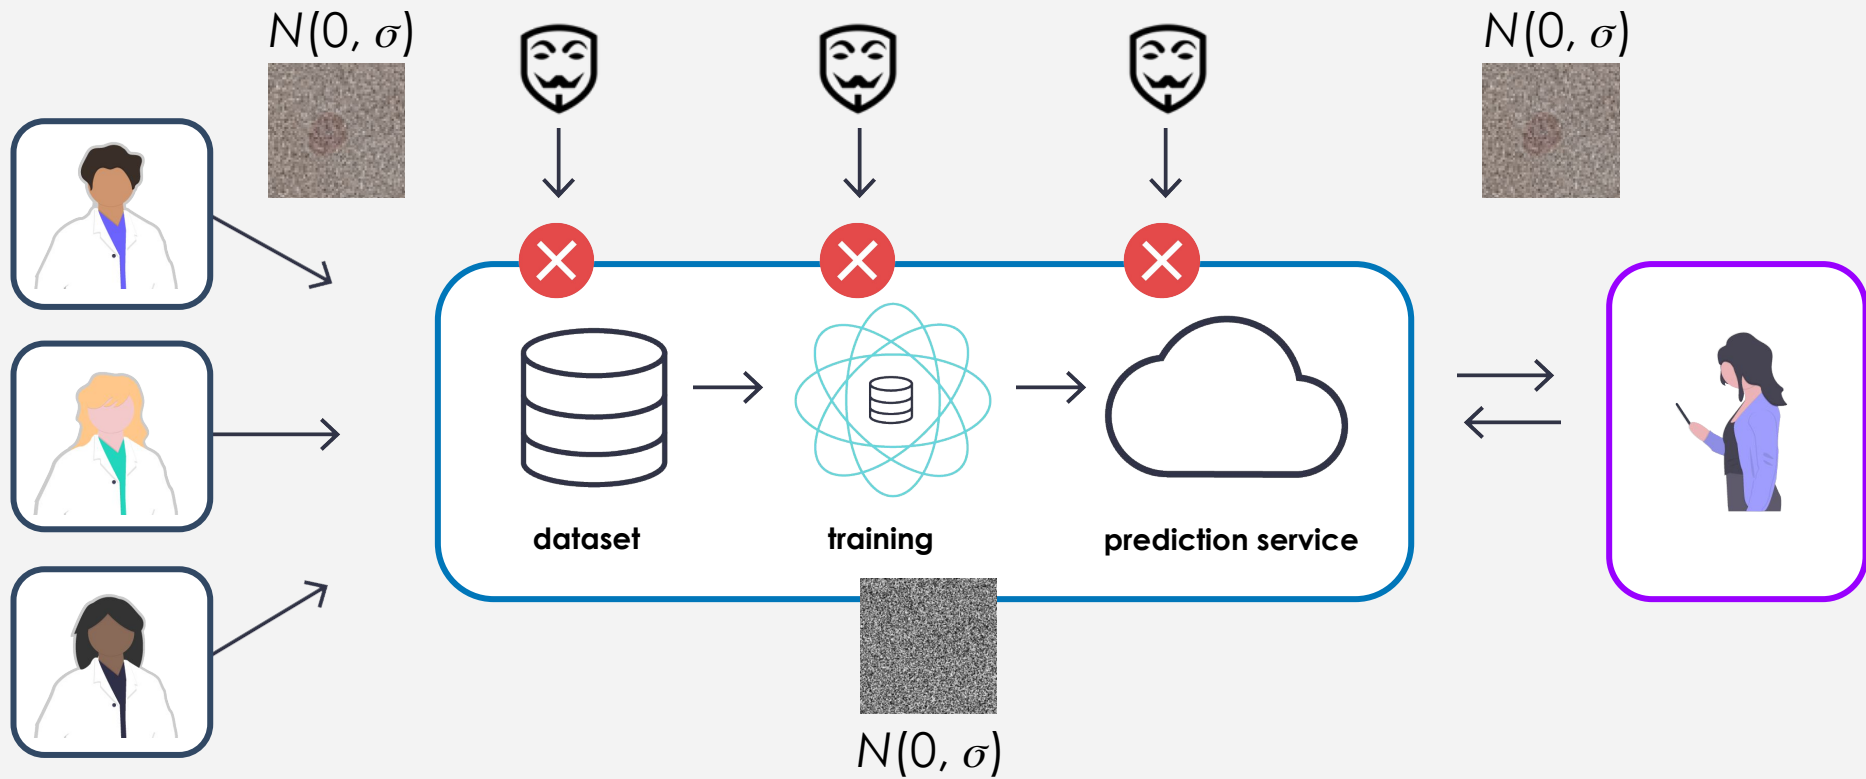
risk management
(store and process)

incentive
(exposure and trust)

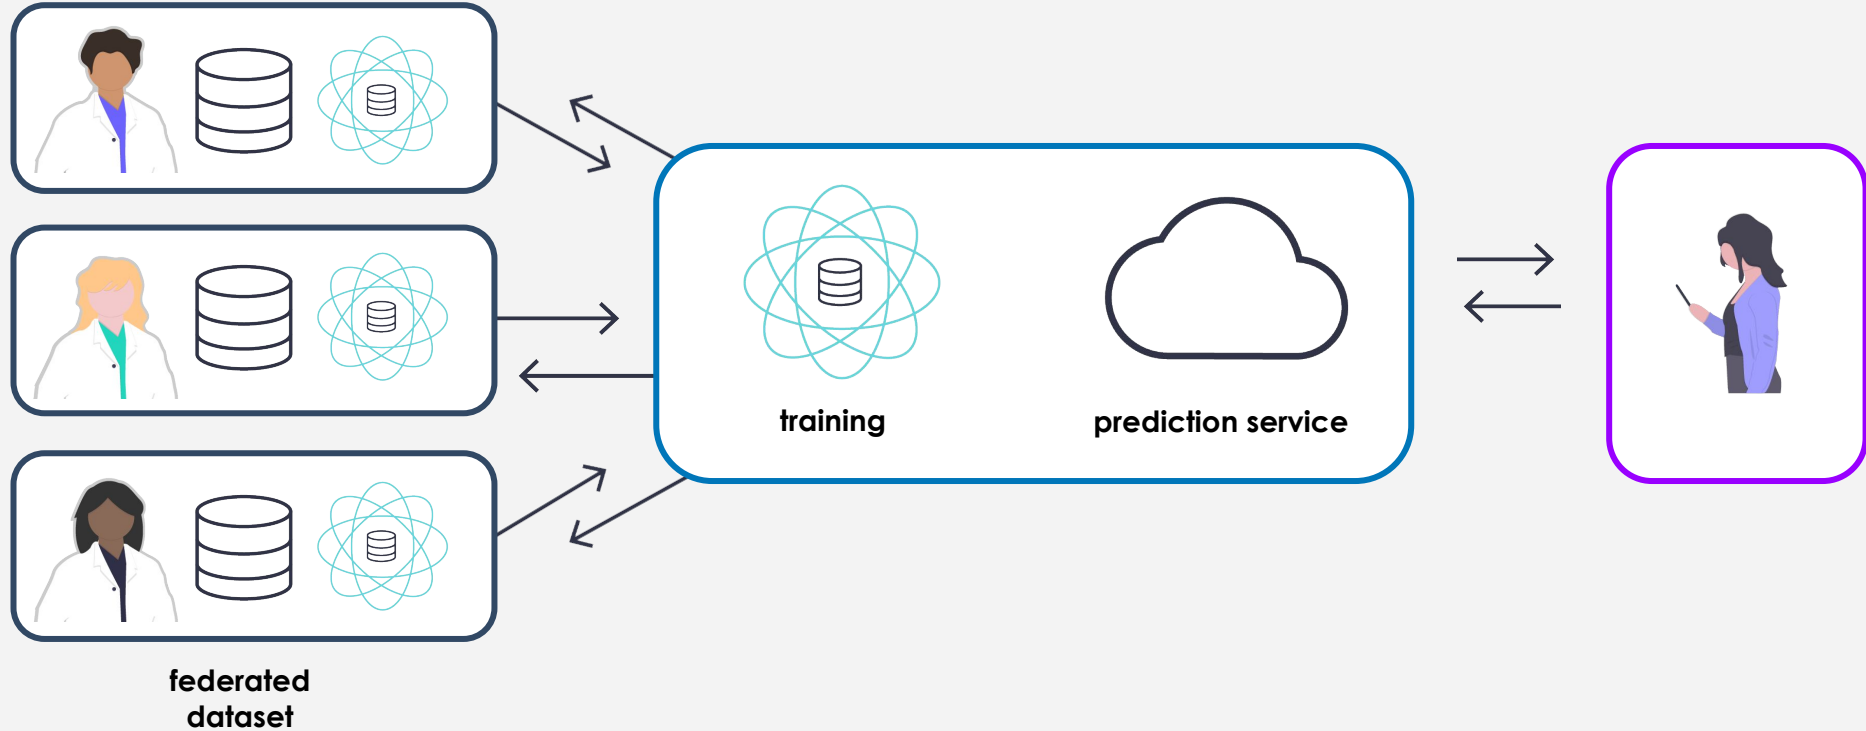


Leakage (model asset
and training data)

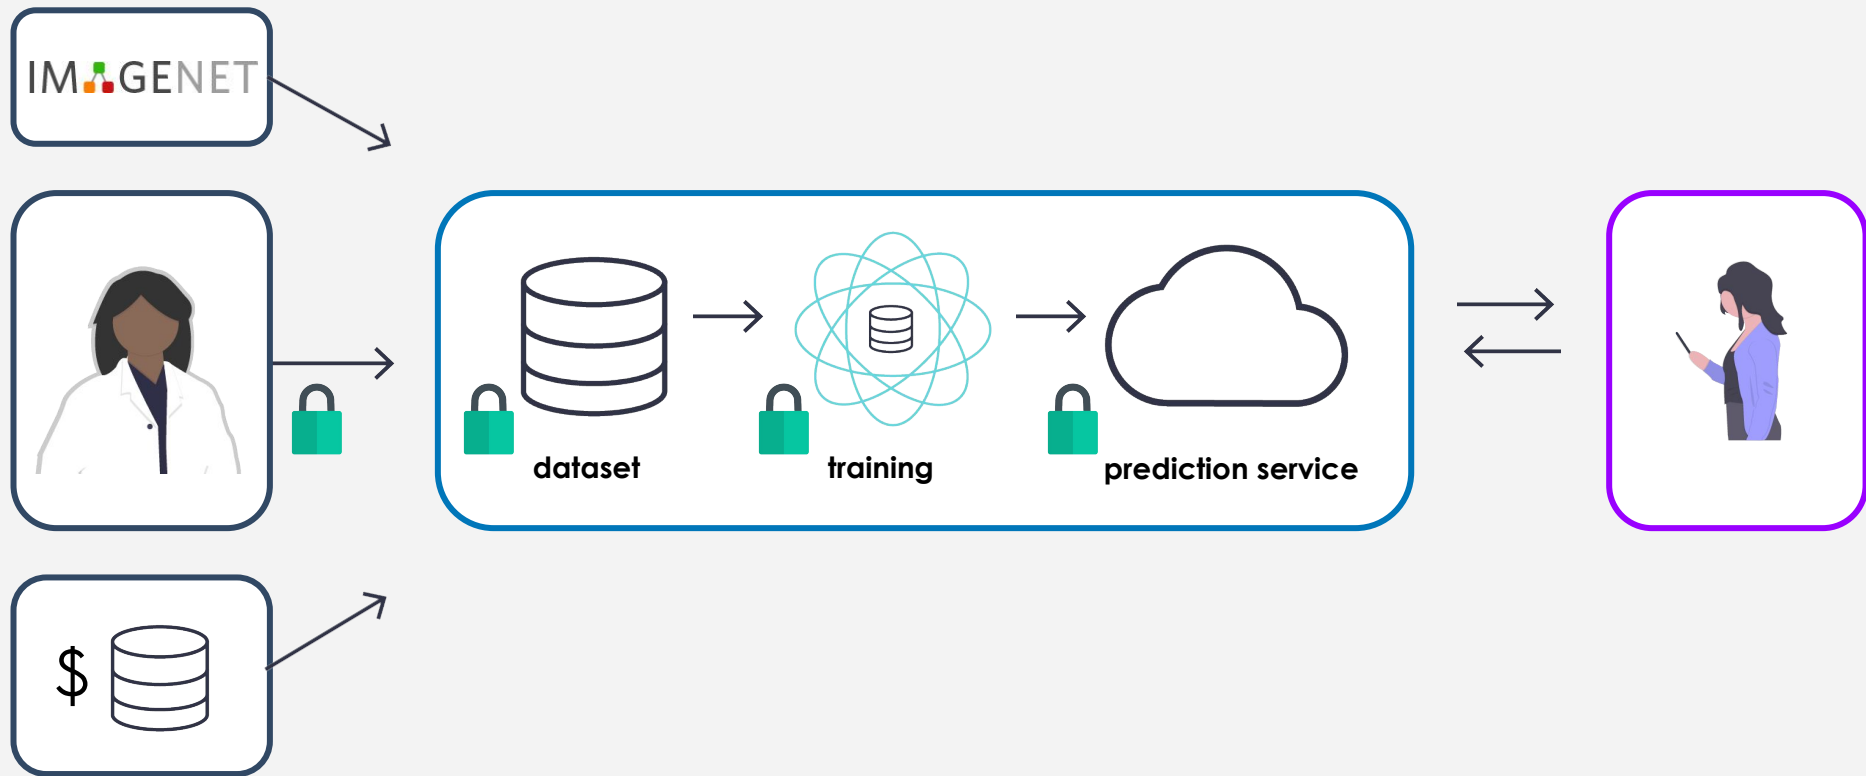
Sanitization (differential privacy)



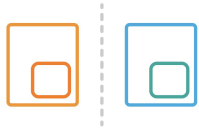
On-Device (federated learning)



Encryption



Privacy Preserving Machine Learning Technologies



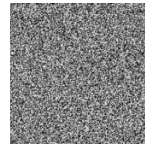
**Remote
Execution**



Federated Learning



**Secure
Computation**



**Differential
Privacy**

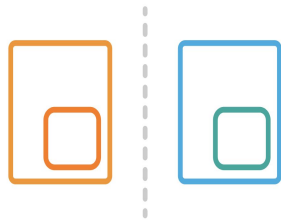
Use Case

Opportunity:

- You want to train a model on sensitive patients' data.

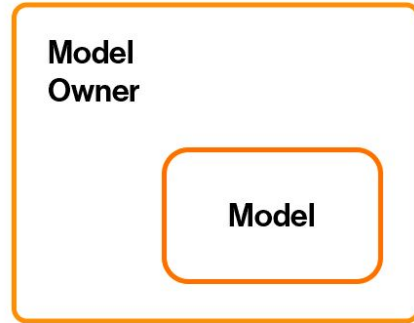
Problem:

- The data is **owned by one hospital**
- It **can't be accessed** directly for liability reasons



**Remote
Execution**

Remote Execution



Why Not Use TensorFlow with Cluster Configuration?

Remote Execution Requirements:

- Access control
- Logging
- Policy
- Audit
- Environment agnostic

Challenges:

- Difficult to control protocol/communication
- Difficult to adapt TensorFlow distributed protocol for other environments (e.g internet)

PySyft

Secure, private Deep Learning library - decouples private data from model training

- Remote execution
- Federated learning
- Differential Privacy
- Multi Party Computation



PySyft - Remote Execution



```
import tensorflow as tf
import syft as sy

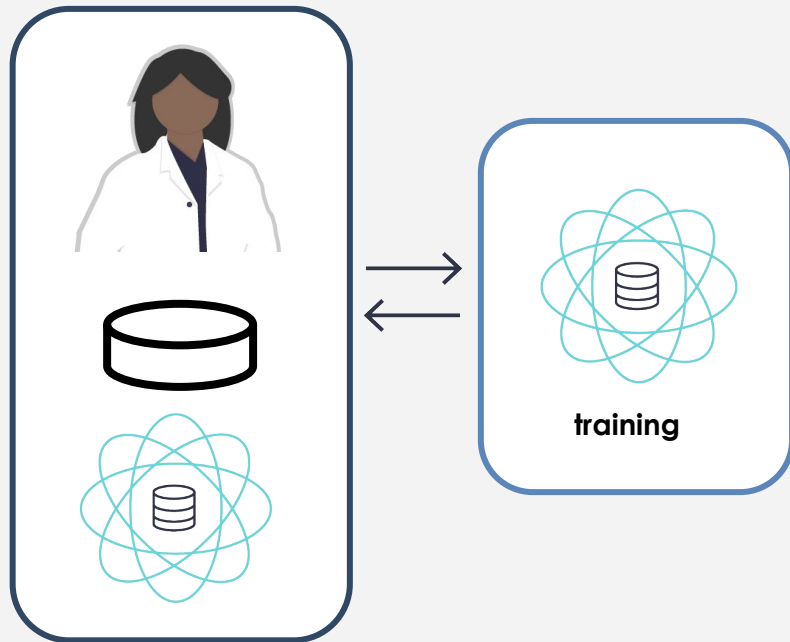
hook = sy.TensorFlowHook(tf)
remote = sy.VirtualWorker(hook, id="remote")

x = tf.constant([1, 2, 3, 4])
y = tf.constant([5, 6, 7, 8])

x_ptr = x.send(remote)
y_ptr = y.send(remote)

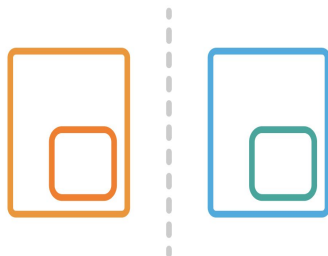
z_ptr = x_ptr + y_ptr

z_ptr.get()
```

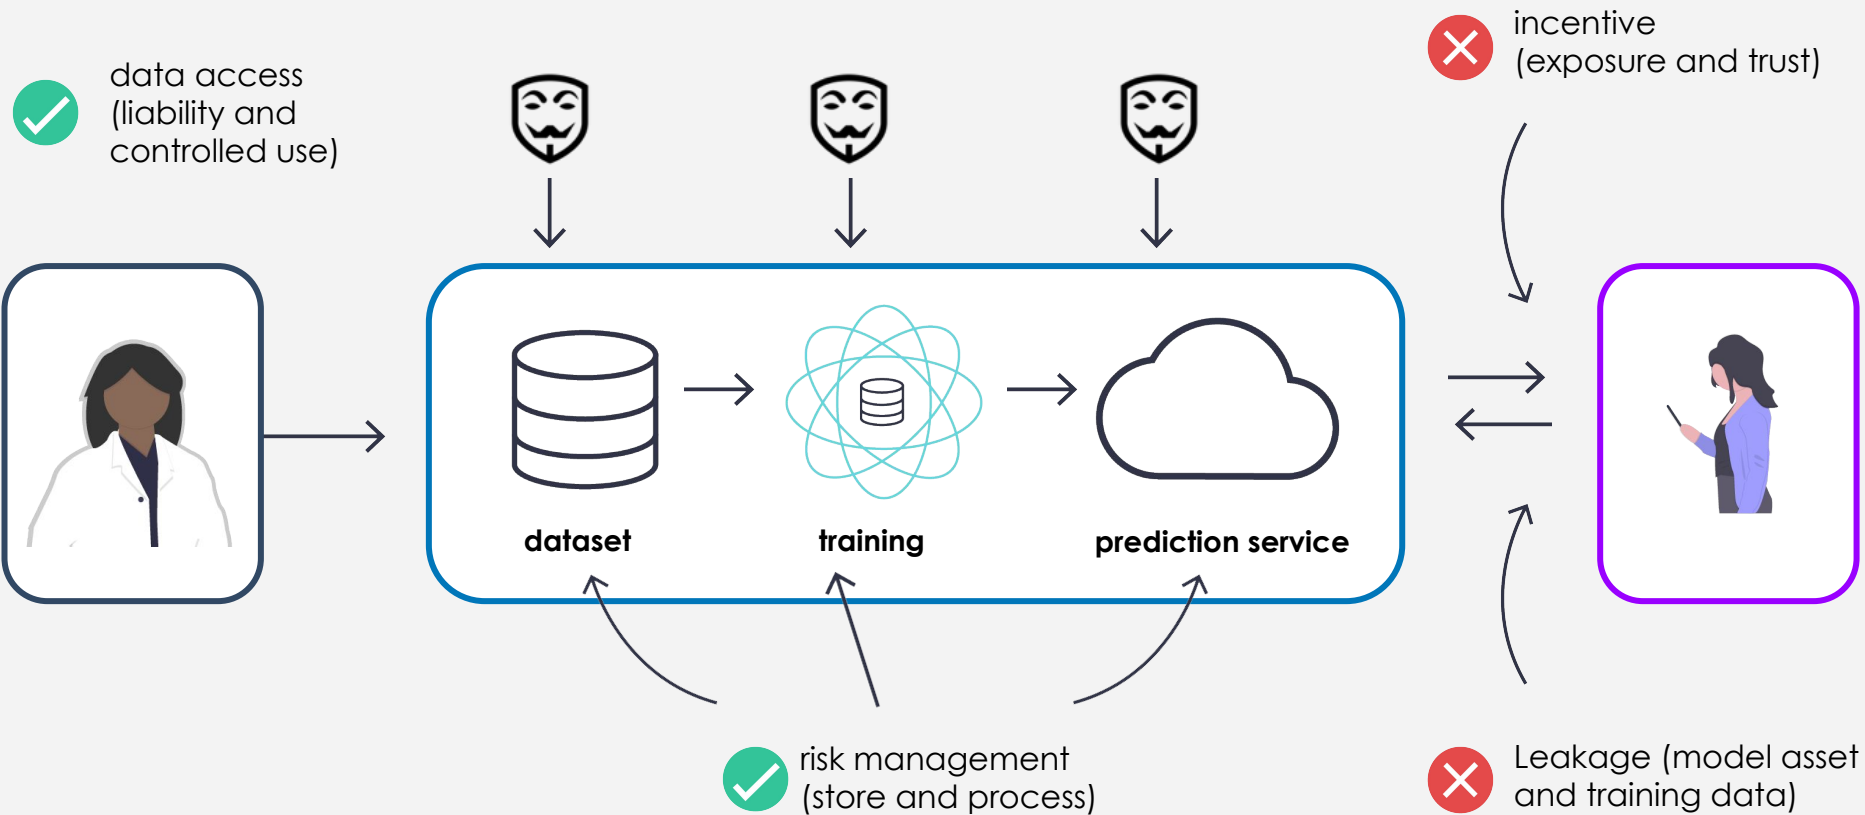


Remote Execution Tutorial

<https://github.com/dropoutlabs/tf-world-tutorial/tree/master/remote-execution>



Remote Execution



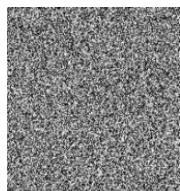
Privacy Leakage With Remote Execution

Problem:

- The model can memorize sensitive data from the training dataset
- Exposed to membership attack and model inversion

Solution:

- Train deep learning model with **Differential Privacy**



Differential Privacy

Differential Privacy

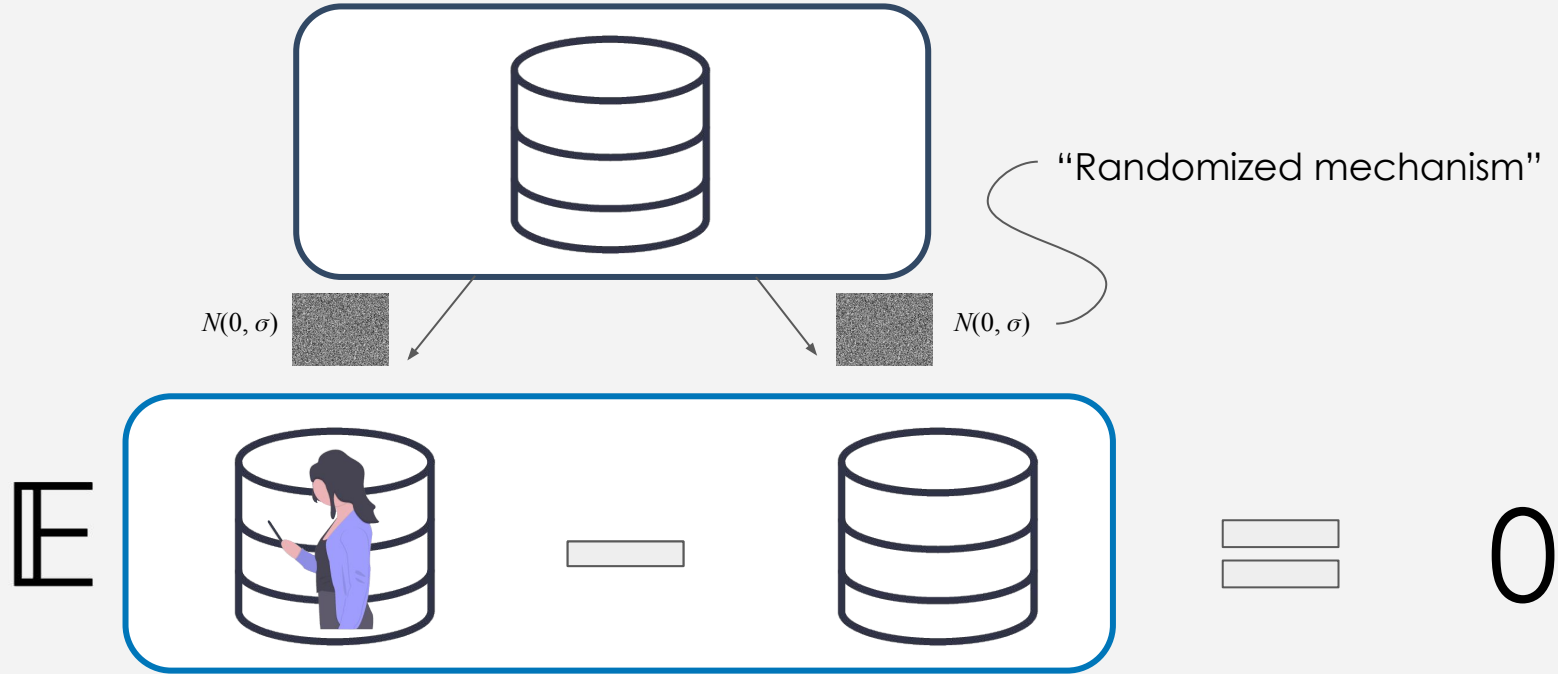
Definition 1. A randomized mechanism $\mathcal{M}: \mathcal{D} \rightarrow \mathcal{R}$ with domain \mathcal{D} and range \mathcal{R} satisfies (ϵ, δ) -differential privacy if for any two adjacent inputs $d, d' \in \mathcal{D}$ and for any subset of outputs $S \subseteq \mathcal{R}$ it holds that

$$\Pr[\mathcal{M}(d) \in S] \leq e^\epsilon \Pr[\mathcal{M}(d') \in S] + \delta.$$

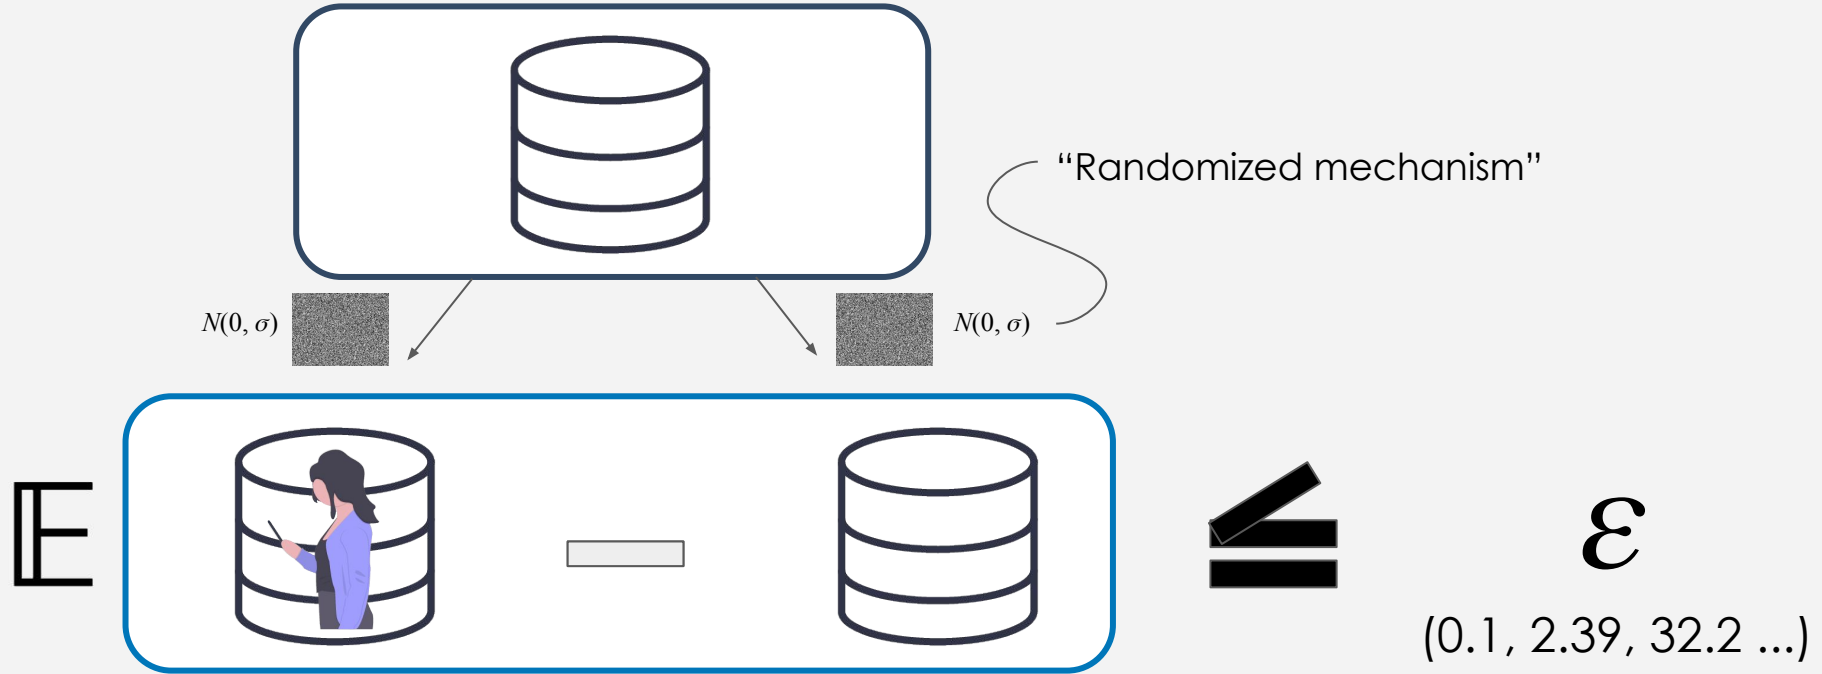
Without Differential Privacy



With Differential Privacy



With Differential Privacy



Privacy vs. Utility

$$D = \{1, 2, 3, 3, 4, n\}$$

$$D' = \{1, 2, 3, 3, 4\}$$

$$f(n) = \text{Avg}(D) - \text{Avg}(D') < \varepsilon$$

$$n \rightarrow \infty ?$$

$$z \sim N(0, \sigma)$$

noise

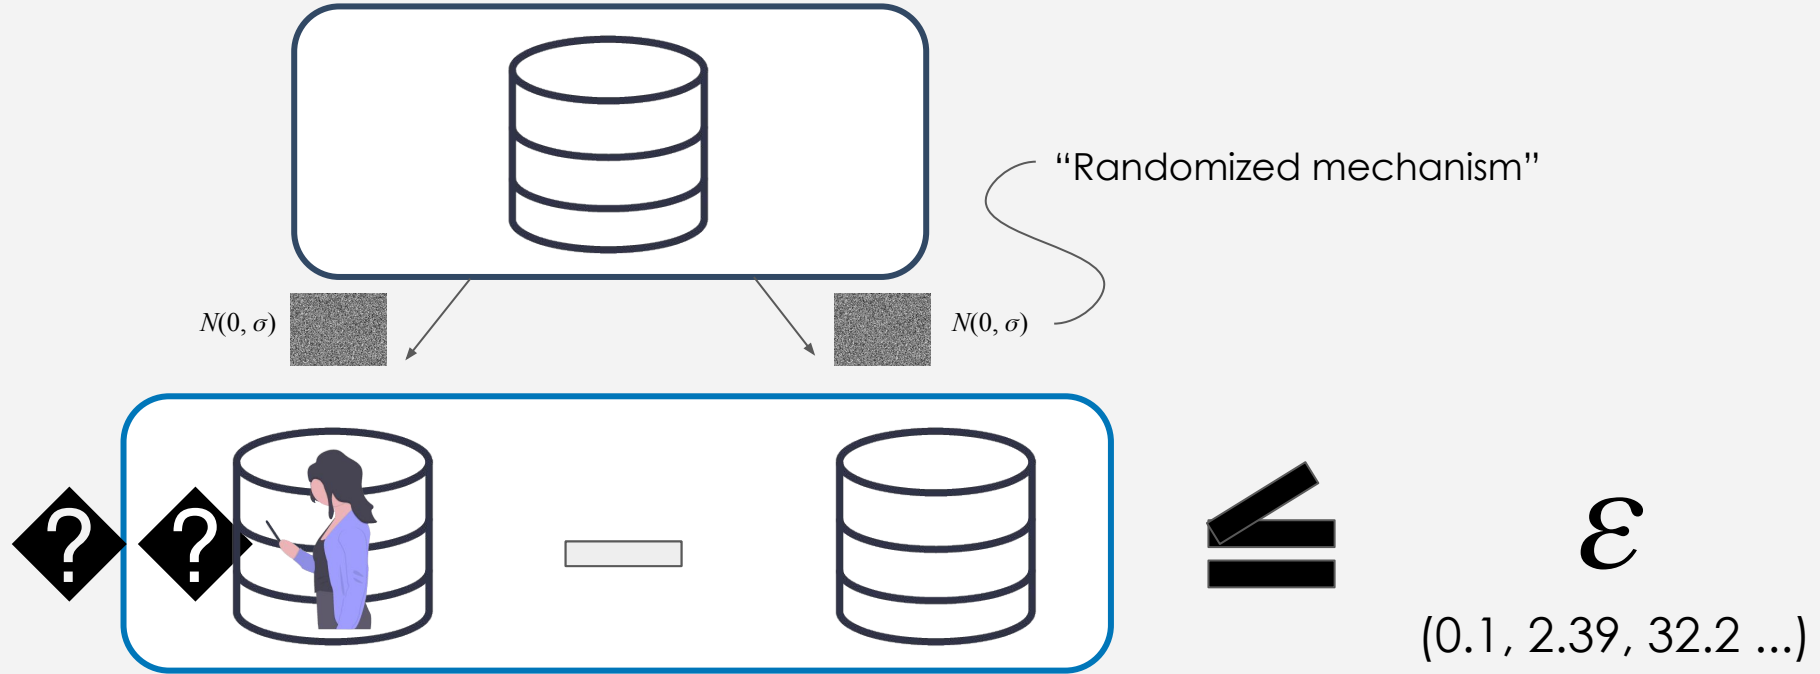
"Randomized mechanism" that
is (ε, δ) -differentially private.

$$f(n) + z \rightarrow \infty$$

$$f(n) \rightarrow \infty - z < \varepsilon$$

$$z \rightarrow \infty$$

With Differential Privacy



Privacy vs. Utility

- Provide optimizers for training models with differential privacy (e.g DP-SGD)
- Ensures that no single such example has any influence, by itself, due to the added noise

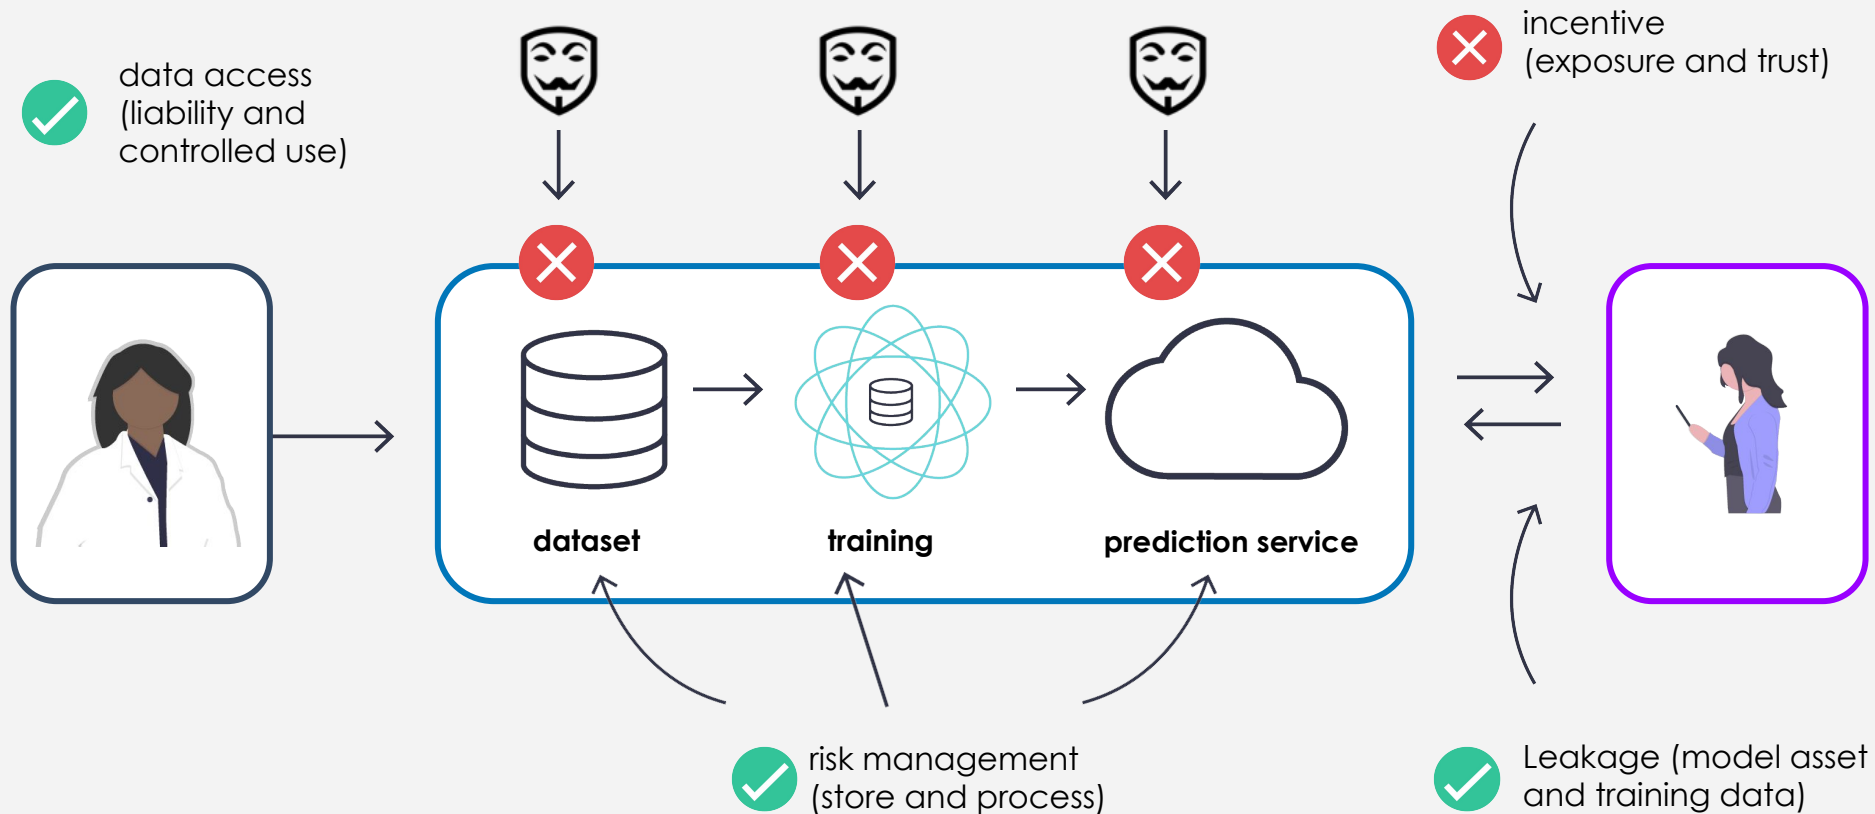
```
optimizer = DPGradientDescentGaussianOptimizer(  
    l2_norm_clip,  
    noise_multiplier,  
    FLAGS.microbatches,  
    FLAGS.learning_rate)  
  
loss = tf.keras.losses.CategoricalCrossentropy(  
    from_logits=True, reduction=tf.losses.Reduction.NONE)  
  
model.compile(optimizer=optimizer, loss=loss, metrics=['accuracy'])
```

- TF Privacy is on PySyft-TensorFlow's roadmap

Further Reading About TensorFlow Privacy

- [Introducing TensorFlow Privacy: Learning with Differential Privacy for Training Data](#)
- [Cleverhans-blog: Privacy and machine learning: two unexpected allies?](#)
- [DPSGD: “Deep Learning with Differential Privacy” Abadi et al.](#)
- [PATE: “Semi-supervised Knowledge Transfer for Deep Learning from Private Training Data” Papernot et al.](#)

Remote Execution + Differential Privacy



Open Questions

Differential privacy does not protect from inferential privacy loss.

If other people give up their data to train models, data about you can be easily/cheaply predicted as long as it belongs to the same distribution -- the privacy loss is contagious.

We currently don't have any good technical solutions to this problem.

PySyft TensorFlow Roadmap

- Federated Learning Support
- Encrypted deep learning: integration with TF Encrypted
- Differential Privacy: integration with TF Privacy
- Add `tf.data` support



PySyft

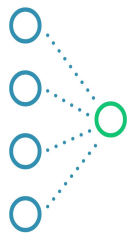
Use Case

Opportunity:

- You want to train a model on sensitive patients' data.

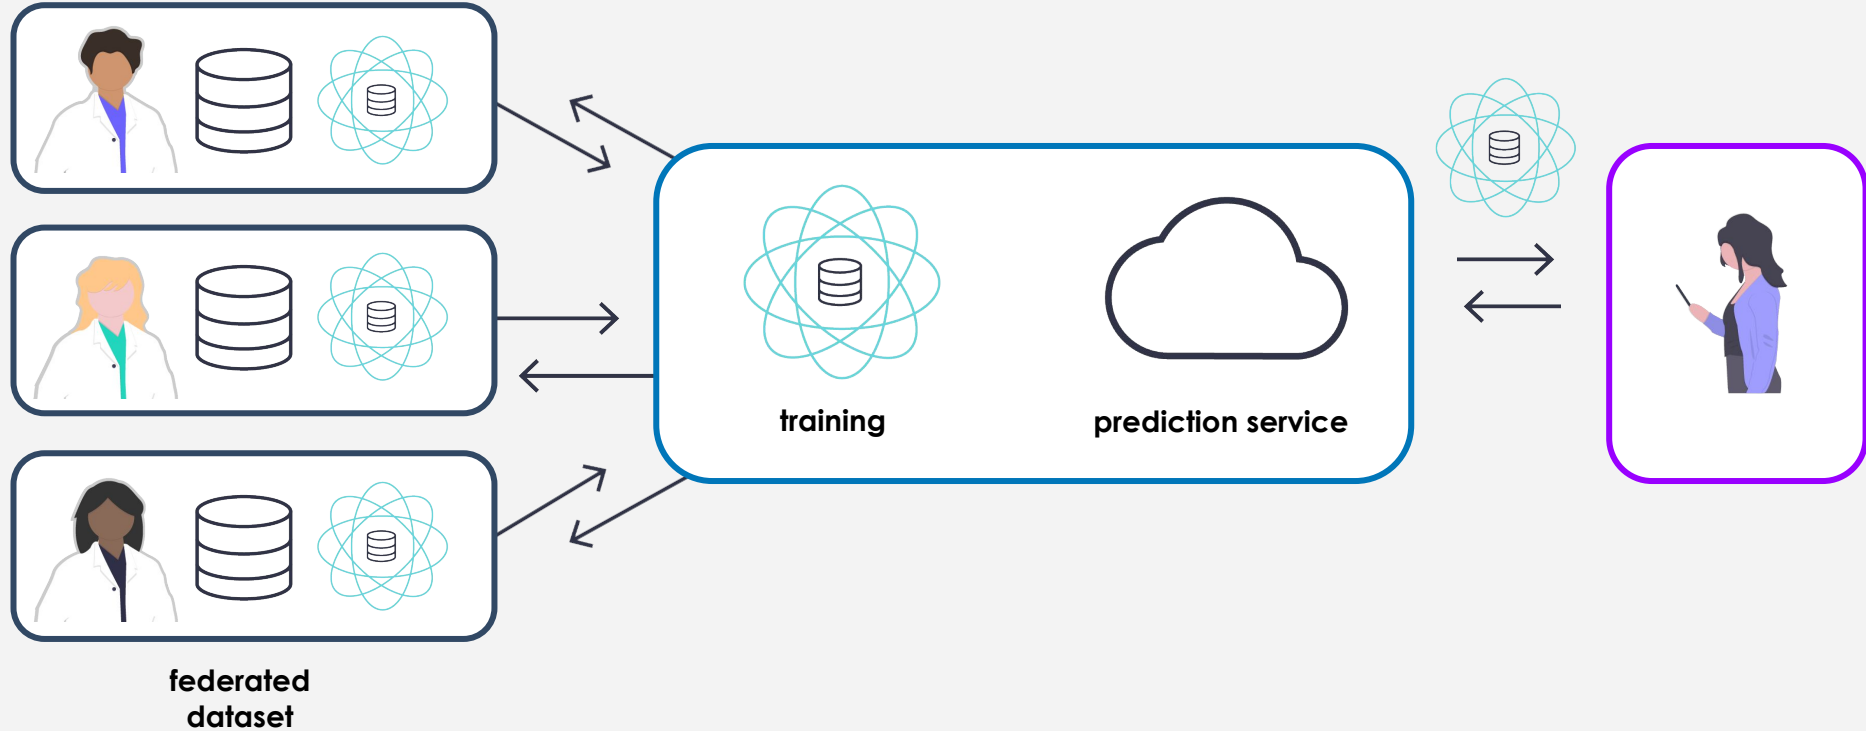
Problems:

- The data is **siloed** across **several hospitals**.
- **The data can't be accessed** directly for liability reasons.



Federated Learning

Federated Learning



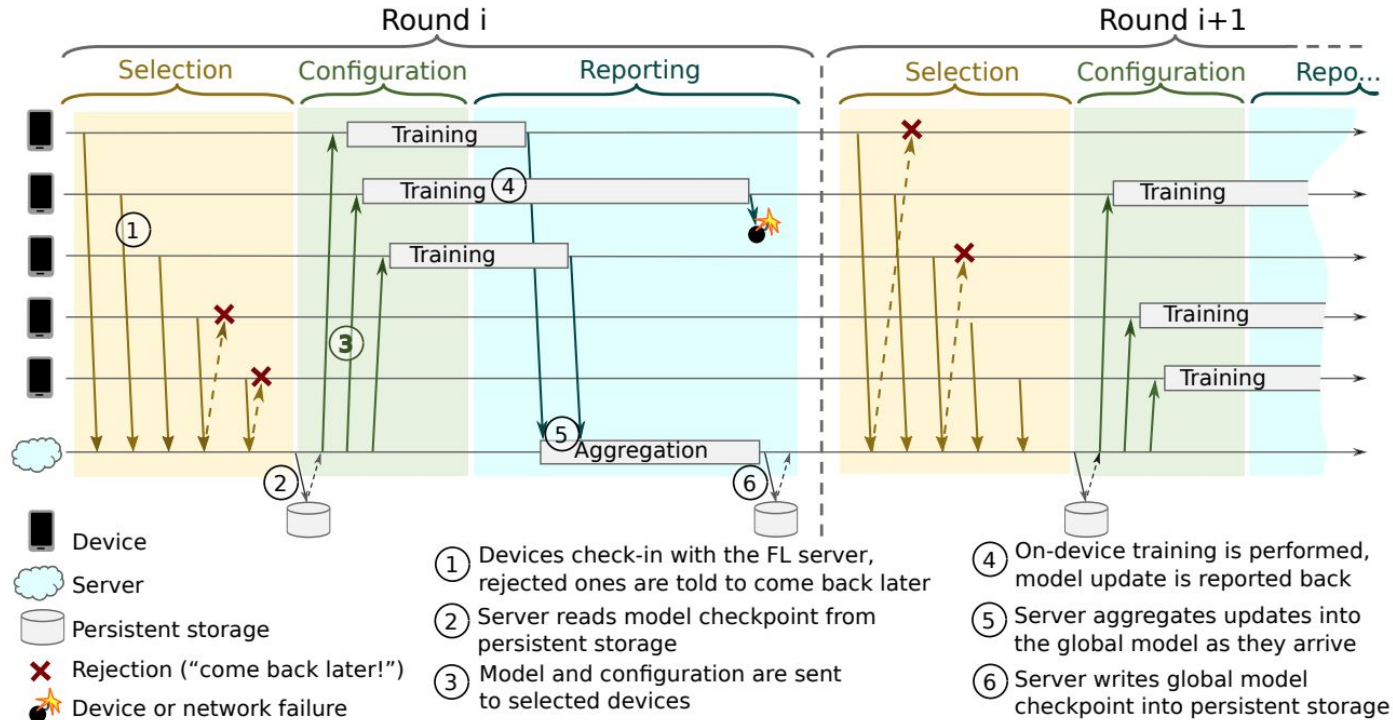


Figure 1: Federated Learning Protocol

[Bonawitz et al. 2019](#)

[TensorFlow Federated](#) on GitHub (Google AI -- Federated Learning Team)



TFEncrypted

TF Encrypted

A Framework for Deep Learning on Encrypted Data

The benefits:

- Usability
- Integration
- Performance
- Extensibility



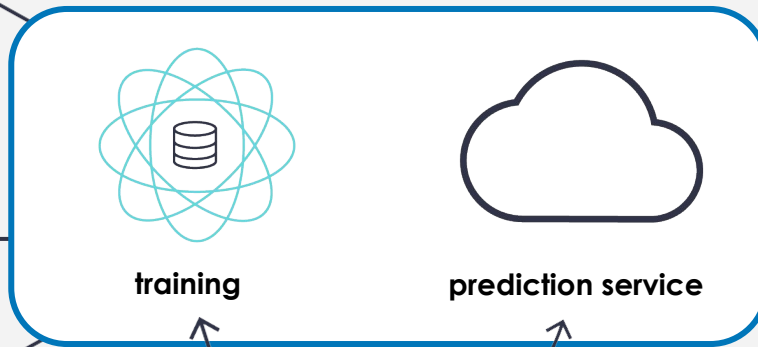
Federated Learning



data access (liability and controlled use)



**federated
dataset**



Risk management
(store and process)



incentive
(exposure and trust)



Leakage (model asset
and training data)

Privacy Leakage With Federated Learning

Problem:

- Can reconstruct data based on gradients updates.

Solution:

- Distribute trust among data owners using **secure aggregation**

Secure Computation

Objective

Evaluate a publicly-known function $f(x, w) = y$ such that x , w , and y are all kept secret¹

Method

Encryptor

Decryptor

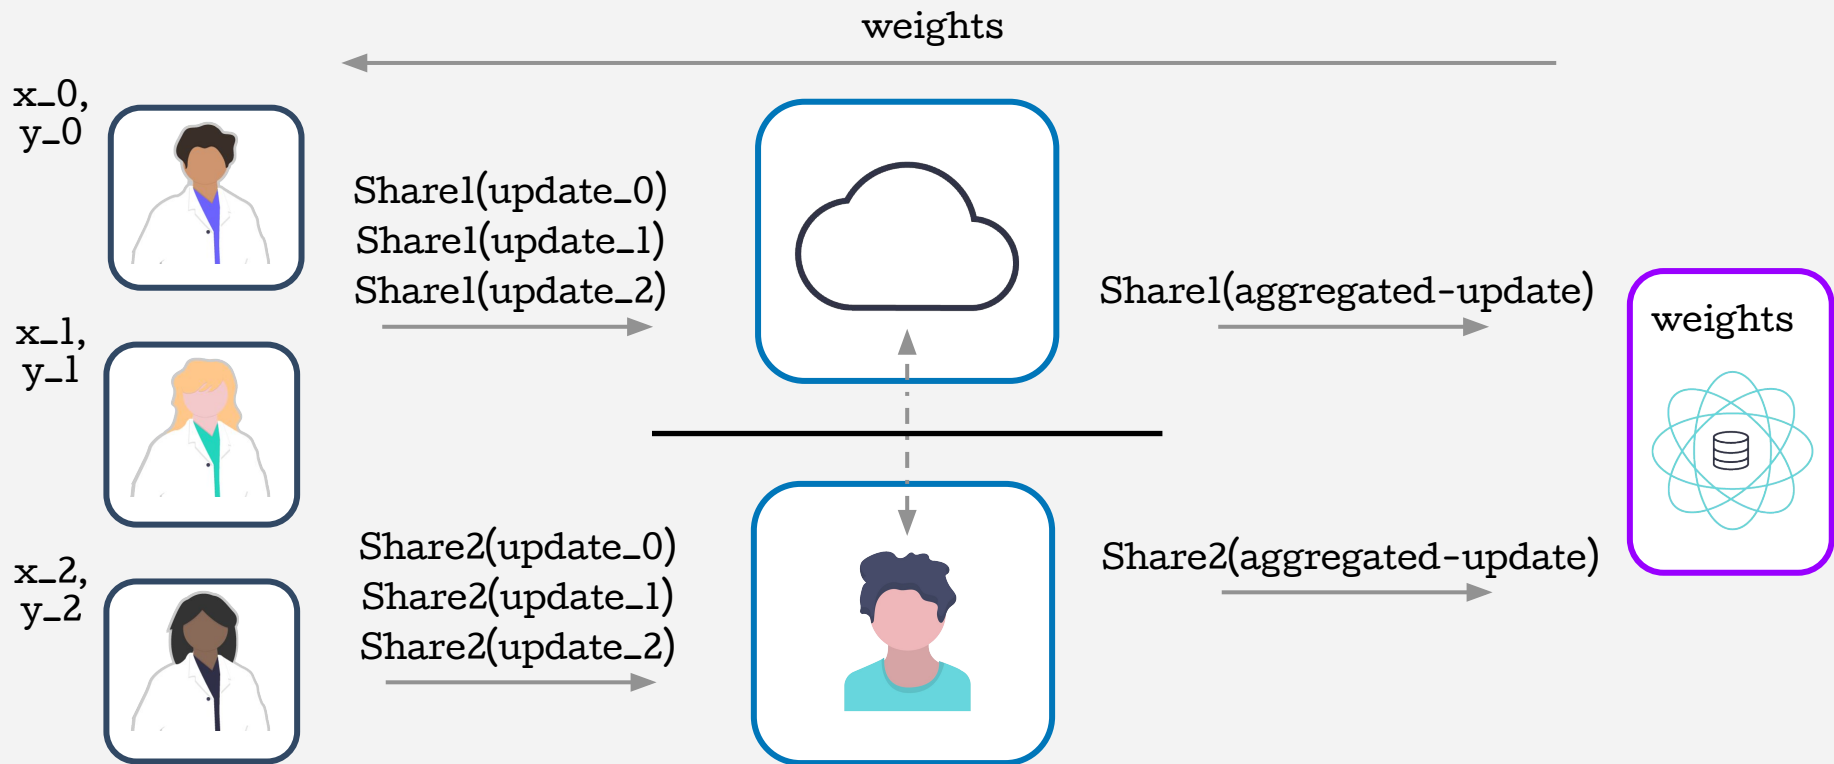
Find a function φ such that φ is homomorphic through f and φ^{-1} is crypto-hard to determine without exact knowledge of some or all inputs.

$$f(\varphi(x), \varphi(w)) = \varphi(f(x, w))$$

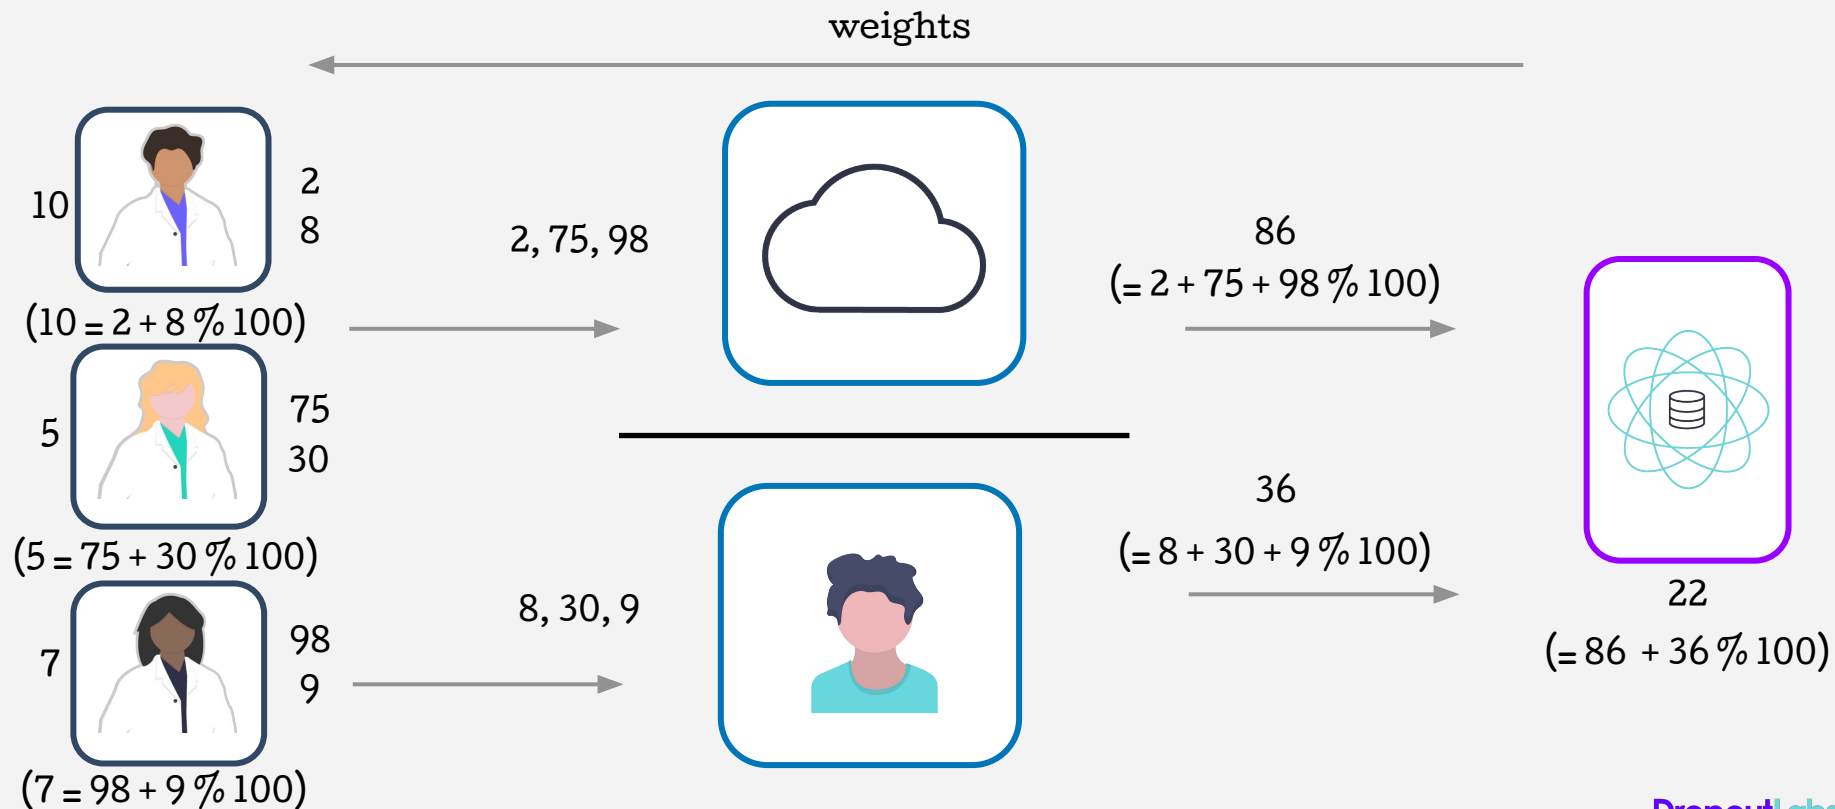
$$(\varphi^{-1} \circ f)(\varphi(x), \varphi(w)) = f(x, w)$$

¹: cryptographically-hard to learn for relevant parties

Participants

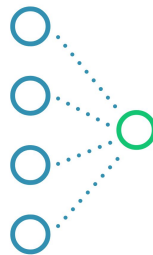


Participants



Secure Federated Learning Tutorial

<https://github.com/dropoutlabs/tf-world-tutorial/tree/master/federated-learning>



Algorithm 1 Reptile (serial version)

Initialize ϕ , the vector of initial parameters

for iteration = 1, 2, ... **do**

 Sample task τ , corresponding to loss L_τ on weight vectors $\tilde{\phi}$

 Compute $\tilde{\phi} = U_\tau^k(\phi)$, denoting k steps of SGD or Adam

 Update $\phi \leftarrow \phi + \epsilon(\tilde{\phi} - \phi)$

end for

Reptile: Pseudocode

Reptile model function:

- Sample batch from DataOwner's dataset

- Run k steps of SGD (or whichever optimizer has been specified)

- Return the new weights

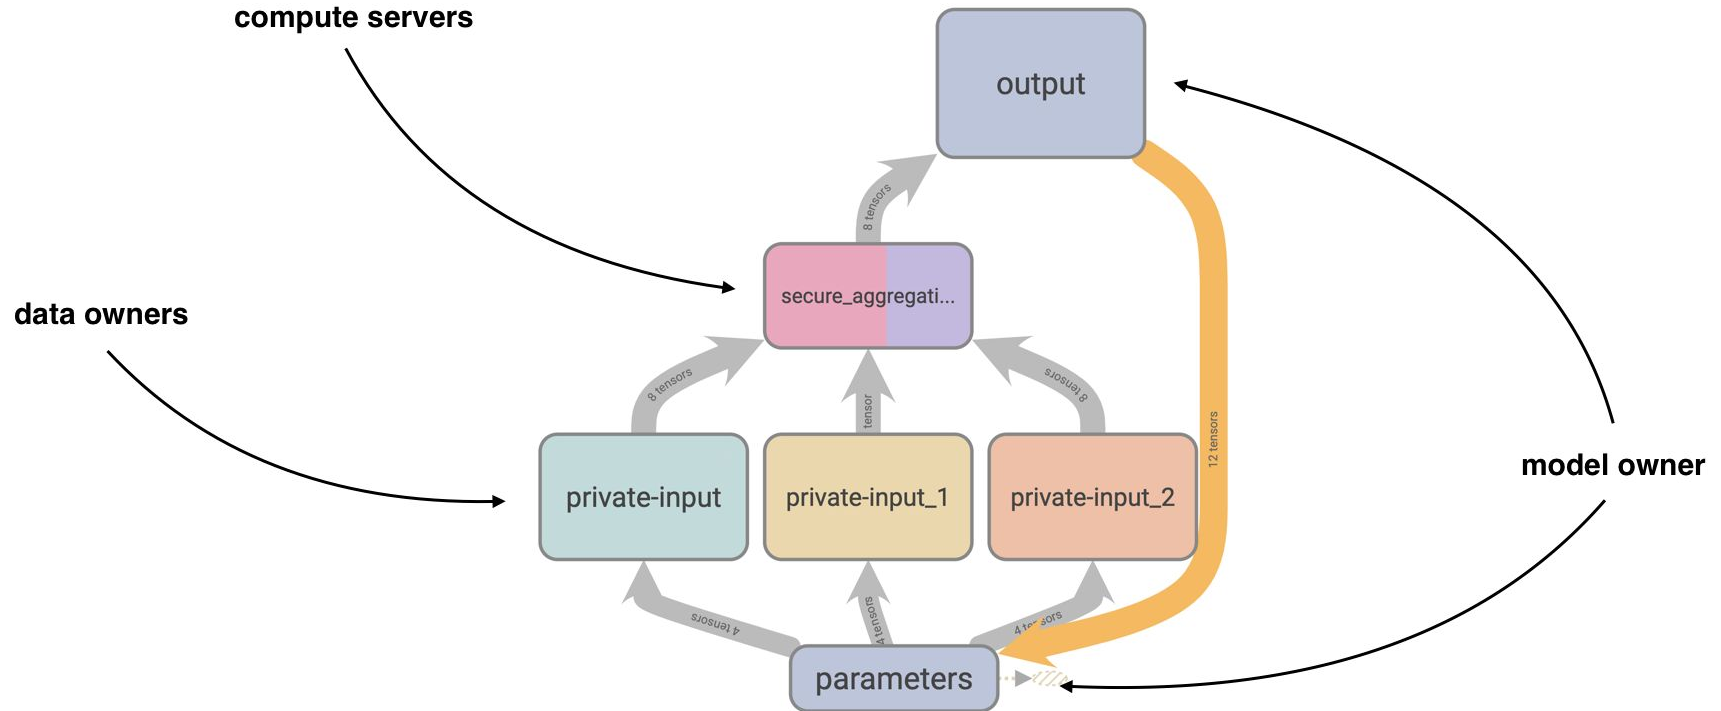
Reptile aggregation function:

- Securely average the model weights

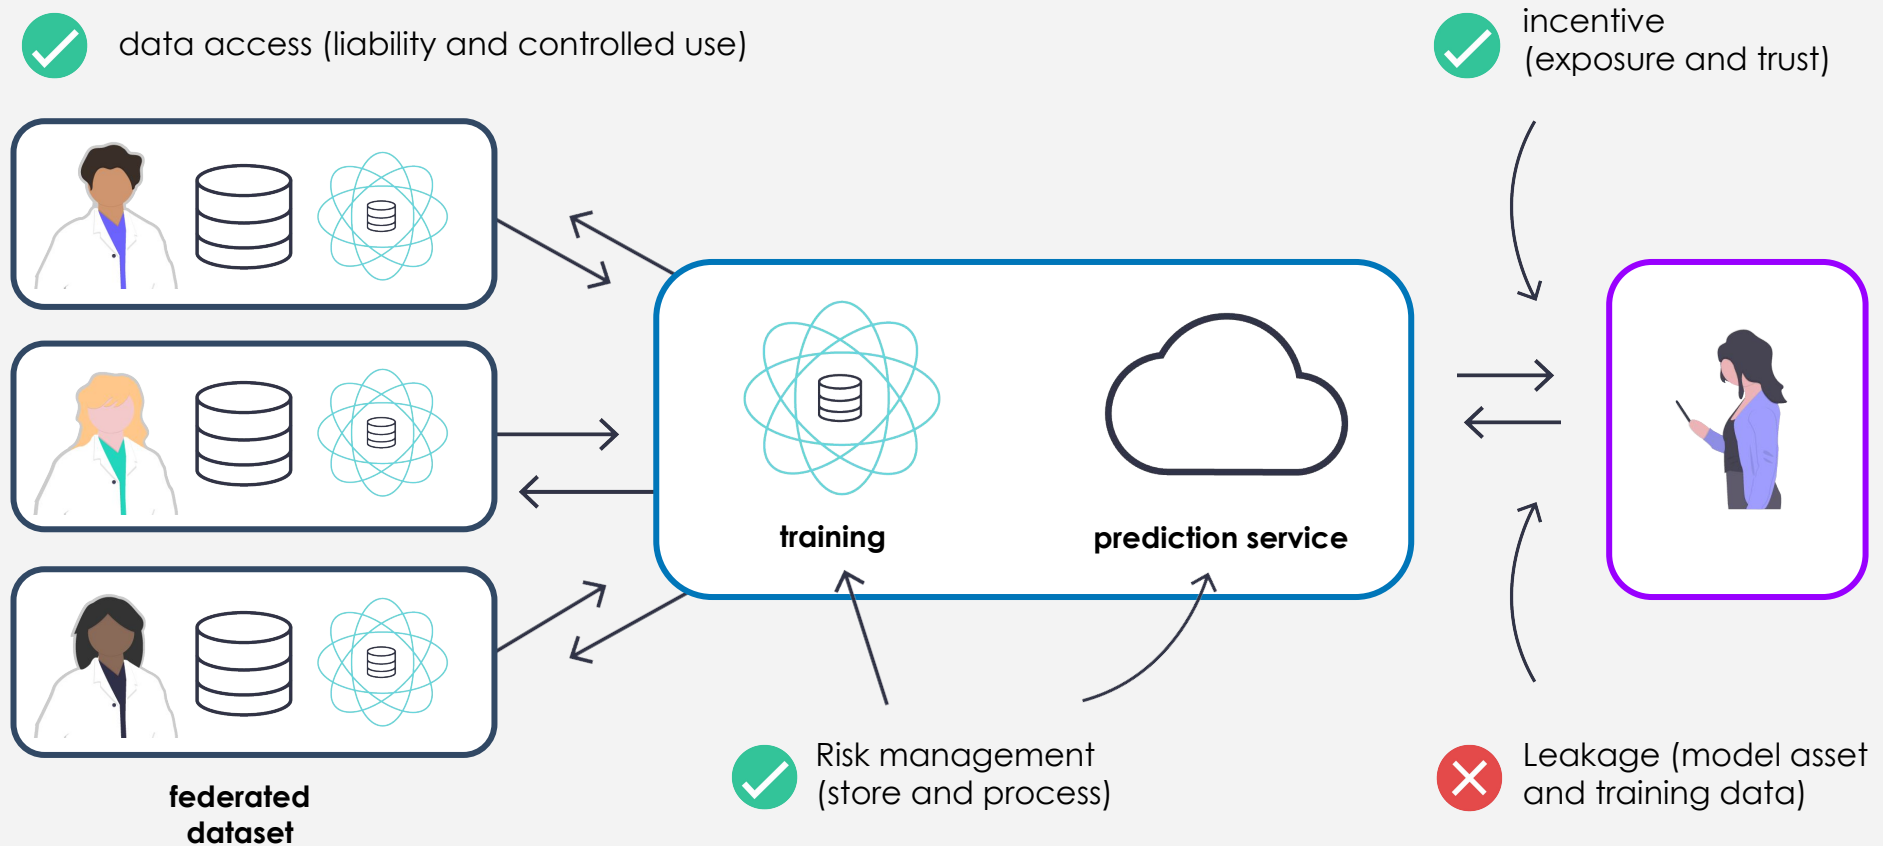
- Subtract newly averaged weights from old weights for approx. gradient

- Update master model with this approximate gradient using SGD.

Overall Computation



Secure Federated Learning



Encrypted Predictions

Third Use Case

Opportunity:

- Provide skin cancer detection to new users.

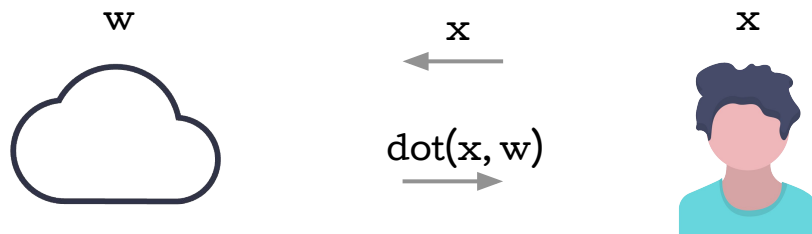
Problems:

- Users won't share their data for privacy reasons.
- Your model is an asset. Ergo, you don't want to send it to their device.



Secure Computation

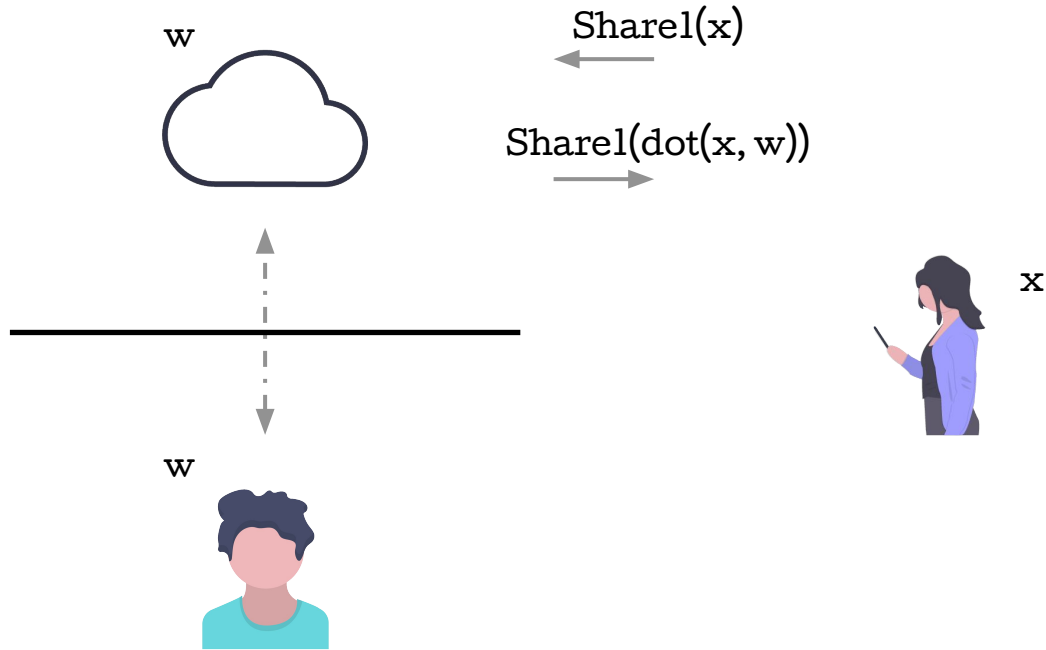
Prediction with Linear Model



$$\text{dot} \left(\begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix}, \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} \right) = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$$



... using Secret Sharing



... using Secret Sharing



Diagram illustrating the computation of the dot product using homomorphic secret sharing:

$$\begin{aligned} \text{dot}(\text{Share}(x_0) \text{ Share}(x_1) \text{ Share}(x_2), \begin{matrix} w_0 \\ w_1 \\ w_2 \end{matrix}) &= \text{Share}(x_0)*w_0 + \text{Share}(x_1)*w_1 + \text{Share}(x_2)*w_2 \\ &= \text{Share}(x_0*w_0) + \text{Share}(x_1*w_1) + \text{Share}(x_2*w_2) \\ &= \text{Share}(x_0*w_0 + x_1*w_1 + x_2*w_2) \end{aligned}$$

Annotations:

- homomorphic secret sharing** (points to the final result)
- public multiplication** (points to the first two terms of the first equation)
- private addition** (points to the second equation)

Secret Sharing

$$\text{Share1}(x, r) = r \bmod m$$

$$\text{Share2}(x, r) = x - r \bmod m$$

$$x = \text{Share1}(x, r) + \text{Share2}(x, r) \bmod m$$

$$m = 10$$

$$\text{Share1}(5, 7) = 7 \bmod 10 = 7$$

$$\text{Share2}(5, 7) = 5 - 7 \bmod 10 = 8$$

$$7 + 8 = 15 = 5 \bmod 10$$



Private Addition with Secret Sharing



x_1

y_1

$$z_1 = x_1 + y_1$$



x_2

y_2

$$z_2 = x_2 + y_2$$

$$x = x_1 + x_2$$

$$y = y_1 + y_2$$

$$\begin{aligned} x + y &= (x_1 + x_2) + (y_1 + y_2) \\ &= (x_1 + y_1) + (x_2 + y_2) \\ &= z_1 + z_2 \end{aligned}$$

Public Multiplication with Secret Sharing



x_1

w

$$z_1 = x_1 * w$$



x_2

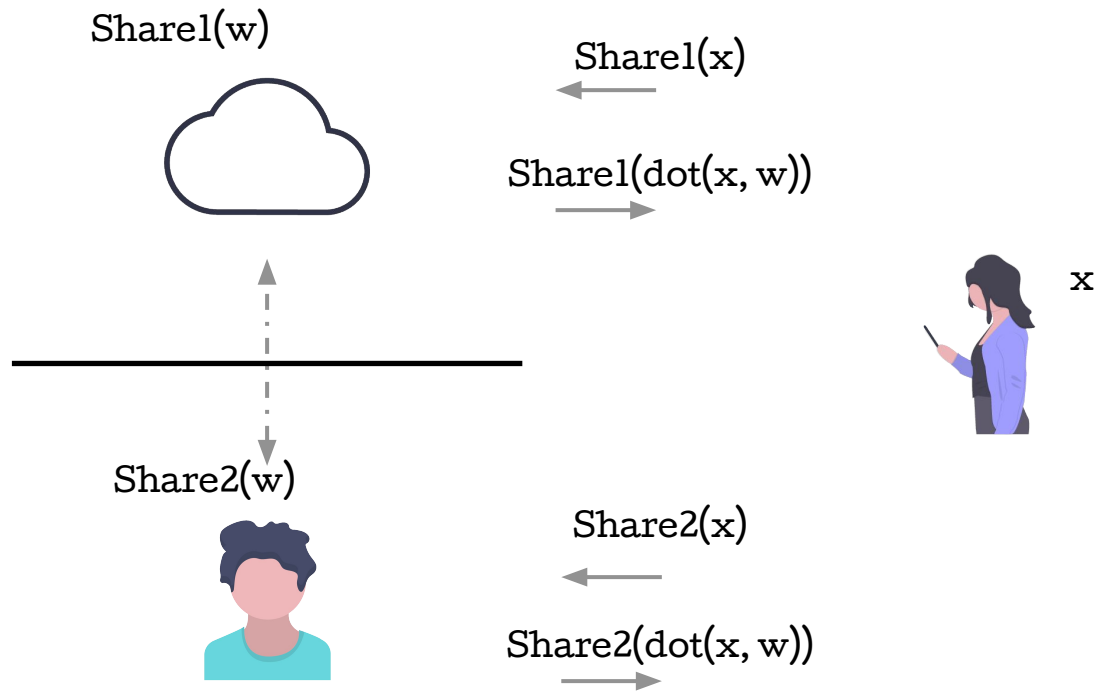
w

$$z_2 = x_2 * w$$

$$x = x_1 + x_2$$

$$\begin{aligned} x * w &= (x_1 + x_2) * w \\ &= (x_1 * w) + (x_2 * w) \\ &= z_1 + z_2 \end{aligned}$$

... using Secret Sharing, with Private Model



... using Secret Sharing, with Private Model

Sharel(**w**)



Sharel(**x**)



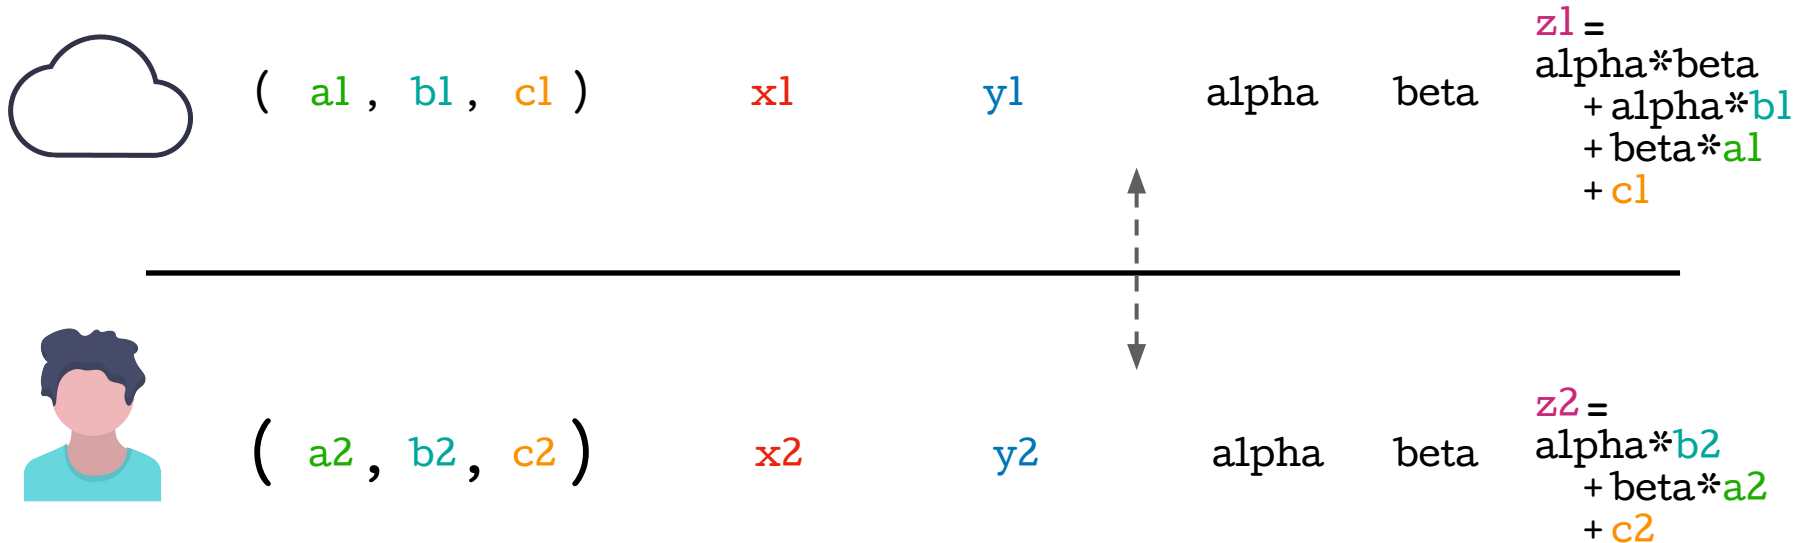
Sharel(dot(**x**, **w**))



private multiplication

$$\begin{aligned} \text{dot}(\text{Sharel}(x_0) \text{ Sharel}(x_1) \text{ Sharel}(x_2), \begin{matrix} \text{Share}(w_0) \\ \text{Share}(w_1) \\ \text{Share}(w_2) \end{matrix}) &= \text{Sharel}(x_0) * \text{Share}(w_0) + \\ &= \text{Sharel}(x_0 * w_0) + \\ &= \text{Sharel}(x_0 * w_0 + x_1 * w_1 + x_2 * w_2) \end{aligned}$$

Private Multiplication with Secret Sharing



$$\begin{aligned} a &= a_1 + a_2 \\ b &= b_1 + b_2 \\ c &= a * b = c_1 + c_2 \end{aligned}$$

$$x = x_1 + x_2 \quad y = y_1 + y_2$$

$$\alpha = x - a$$

$$\beta = y - b$$

$$\begin{aligned} x * y &= \\ &= z_1 + z_2 \end{aligned}$$

Encrypted Predictions Tutorial

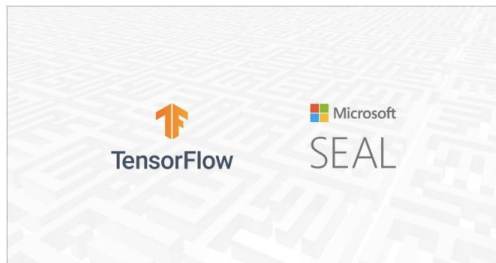
<https://github.com/dropoutlabs/tf-world-tutorial/tree/master/federated-learning>



Characteristics

	Secure Enclaves	Garbled Circuits	Homomorphic Encryption	Secret Shares
Compute	Low	Medium	Very High	Low
Network	Low	Medium	Low	High
Operations	Any	Boolean	Arithmetic	Arithmetic
Runs on	Hardware	Software	Software	Software
Parties	2+	2+	2+	3+

TF Seal: Homomorphic Encryption coming soon to TF Encrypted!



 Morten Dahl and 21 others

Bridging Microsoft SEAL into TensorFlow

The road to machine learning with homomorphic encryption



Justin Patriquin

Aug 8 · 6 min read

GitHub:

<https://github.com/tf-encrypted/tf-seal>

Medium:

<https://medium.com/dropoutlabs/bridging-microsoft-seal-into-tensorflow-b04cc2761ad4>

TF Trusted: Secure Enclaves coming soon to TF Encrypted!



Ainsley Sutherland and 21 others

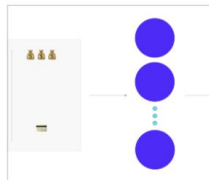
Secure Logistic Regression: MPC vs Enclave Benchmark

With Ben DeCoste.



Justin Patriquin

Feb 6 · 7 min read



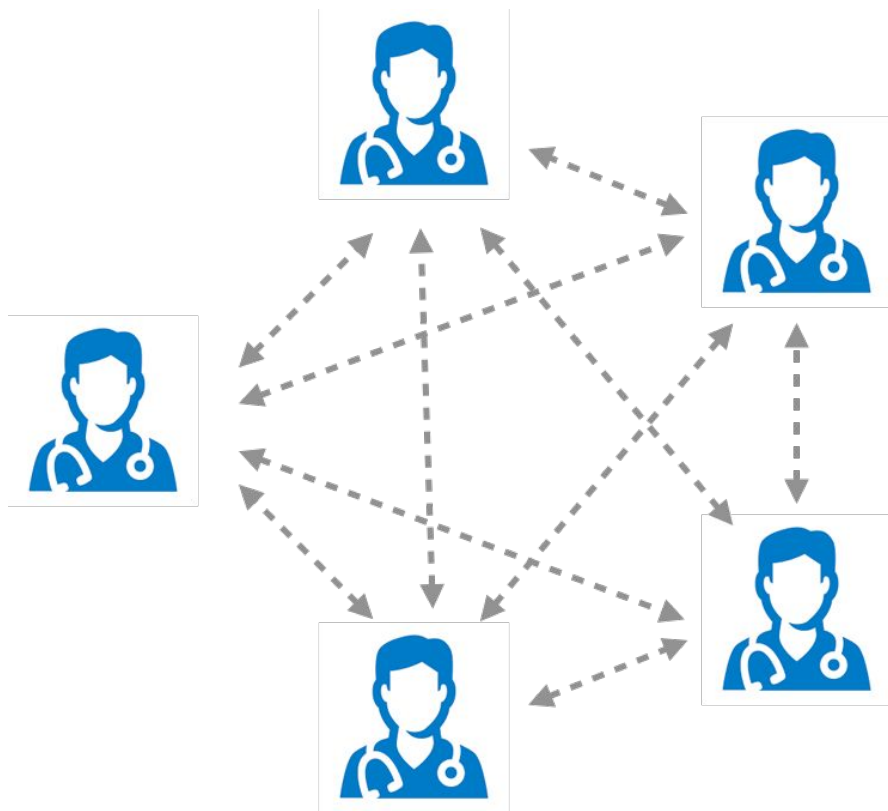
Github:

<https://github.com/dropoutlabs/tf-trusted>

Medium:

<https://medium.com/dropoutlabs/secure-logistic-regression-mpc-vs-enclave-benchmark-65cf5f03a81f>

Multi-party training



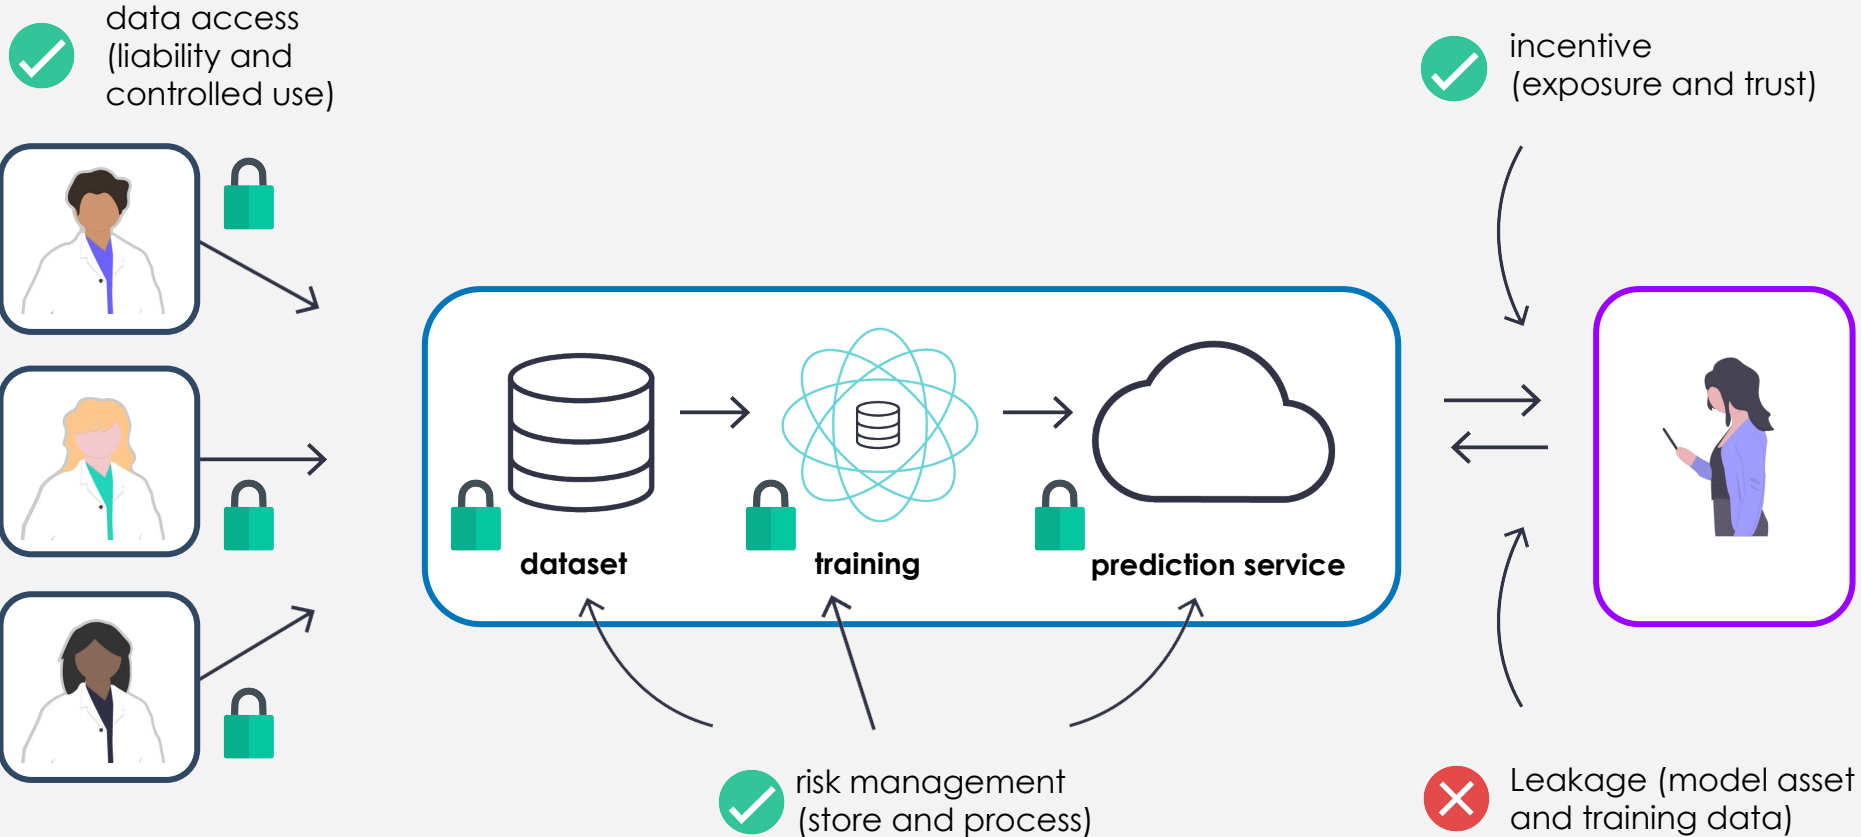
Share1(w)



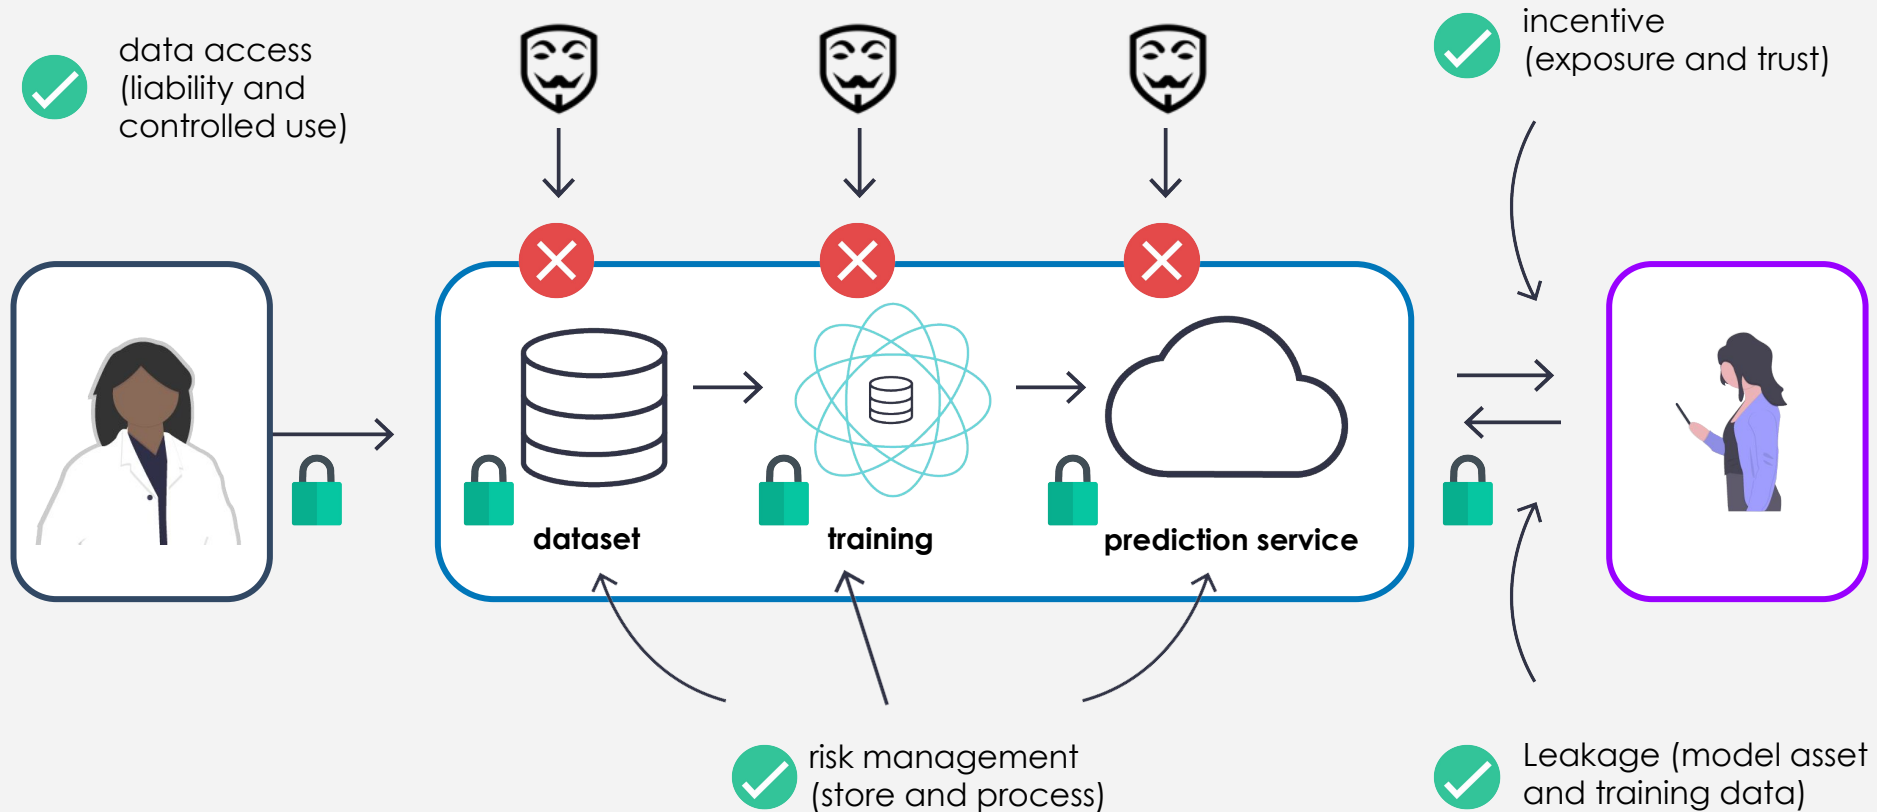
Share2(w)



Secure Computation



Secure Computation + Differential Privacy



Open Questions

Can we make the currently prohibitive performance cost marginal? Or is secure computation fated to only enable use cases that aren't currently possible (e.g. due to regulatory concerns)

Or can we combine MPC with other approaches like federated learning to give it a privacy guarantee? (e.g. secure aggregation from Google Federated Learning team)

Similar to what people are trying to do now with differential privacy, are there alternative definitions of privacy/security that would lead to better performance?

PrivacyAI

www.privacy.ai

A Platform for Secure, Privacy-Preserving Machine Learning

Build **better models** by working with **more valuable data**

[Request Demo](#)



Self-serve access to
protected data



Integrated with your
existing workflow



Powered by modern
privacy techniques



Governed by security
policies and permissions